

```
!pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.51.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.18.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.30.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.31.2)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.21.1)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.3)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0->transformers) (2024.11.6)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0->transformers) (4.6.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.4.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2025.4.26)
```

```
from transformers import pipeline
```

```
a = pipeline("sentiment-analysis")
```

```
No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-2-english and revision 714eb0f (https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english).
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 629/629 [00:00<00:00, 39.5kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, please install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100% 268M/268M [00:01<00:00, 237MB/s]

tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 2.99kB/s]

vocab.txt: 100% 232k/232k [00:00<00:00, 4.29MB/s]

Device set to use cpu
```

```
text = input("Enter the text:")
print(a(text))
```

```
Enter the text:i love this product.its amazing and works perfectly
[{'label': 'POSITIVE', 'score': 0.9998792409896851}]
```

```
gen = pipeline('text-generation')
```

```
No model was supplied, defaulted to openai-community/gpt2 and revision 607a30d (https://huggingface.co/openai-community/gpt2).
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 665/665 [00:00<00:00, 29.9kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, please install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100% 548M/548M [00:02<00:00, 236MB/s]

generation_config.json: 100% 124/124 [00:00<00:00, 9.35kB/s]

tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 1.78kB/s]

vocab.json: 100% 1.04M/1.04M [00:00<00:00, 22.4MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 20.1MB/s]

tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 27.1MB/s]

Device set to use cpu
```

```
text = input("Enter the text:")
t = (gen(text, max_length = 200))
```

```
Enter the text:once upon a time there lived a lion king mufasa
Setting 'pad_token_id' to 'eos_token_id':50256 for open-end generation.
```

```
print(t[0].get('generated_text'))
```

once upon a time there lived a lion king mufasa,' " says the Bovine scholar. The chief prince may have held the title while the lion kin

A second-mentioned historical fact is that in many old accounts of the kings' life, their lionhood came about by the custom of offering

But for it is not true that the king of Norway is a niyasa. The king is a bovine and was once a lion king (although he lived before such

```
n = pipeline('ner')
```

No model was supplied, defaulted to dbmdz/bert-large-cased-finetuned-conll03-english and revision 4c53496 (<https://huggingface.co/dbmdz/bert-large-cased-finetuned-conll03-english>)
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 998/998 [00:00<00:00, 80.4kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100% 1.33G/1.33G [00:10<00:00, 280MB/s]

Some weights of the model checkpoint at dbmdz/bert-large-cased-finetuned-conll03-english were not used when initializing BertForTokenClassification
- This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task or with another architecture
- This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be exactly identical to

tokenizer_config.json: 100% 60.0/60.0 [00:00<00:00, 3.89kB/s]

vocab.txt: 100% 213k/213k [00:00<00:00, 5.33MB/s]

Device set to use cpu

```
text = input("Enter a text")
result = n(text)
result
```

Enter a text: am ahalya from guntur, studying in vignan university

```
[{'entity': 'I-PER',
  'score': np.float32(0.537826),
  'index': 3,
  'word': 'ah',
  'start': 5,
  'end': 7},
 {'entity': 'I-LOC',
  'score': np.float32(0.49098775),
  'index': 7,
  'word': 'gun',
  'start': 17,
  'end': 20},
 {'entity': 'I-LOC',
  'score': np.float32(0.8940649),
  'index': 8,
  'word': 'vign',
  'start': 20,
  'end': 23},
 {'entity': 'I-ORG',
  'score': np.float32(0.5598335),
  'index': 13,
  'word': 'university',
  'start': 38,
  'end': 41}]
```

```
qa = pipeline('question-answering')
```

No model was supplied, defaulted to distilbert/distilbert-base-cased-distilled-squad and revision 564e9b5 (<https://huggingface.co/distilbert/distilbert-base-cased-distilled-squad>)
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 473/473 [00:00<00:00, 38.7kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100% 261M/261M [00:01<00:00, 225MB/s]

tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 3.66kB/s]

vocab.txt: 100% 213k/213k [00:00<00:00, 9.41MB/s]

tokenizer.json: 100% 436k/436k [00:00<00:00, 23.9MB/s]

Device set to use cpu

```
c = '''Bity bought some butter, but the butter was bitter butter.
So, Bitty bought some better butter, to make bitter butter better'''
```

```
q = "What does Bitty bought?"
```

```
result = qa({'question': q, 'context': c})
print(result)
```

```
{'score': 0.193147212266922, 'start': 76, 'end': 94, 'answer': 'some better butter'}
/usr/local/lib/python3.11/dist-packages/transformers/pipelines/question_answering.py:390: FutureWarning: Passing a list of SQuAD example
warnings.warn(
```

```
s = pipeline('summarization')
```

```
No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (https://huggingface.co/sshleifer/distilbart-cnn-12-6).
Using a pipeline without specifying a model name and revision in production is not recommended.
```

```
config.json: 100% 1.80k/1.80k [00:00<00:00, 109kB/s]
pytorch_model.bin: 100% 1.22G/1.22G [00:13<00:00, 43.2MB/s]
model.safetensors: 100% 1.22G/1.22G [00:44<00:00, 205MB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 1.36kB/s]
vocab.json: 100% 899k/899k [00:00<00:00, 1.80MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 4.26MB/s]
Device set to use cpu
```

```
t = '''In the 16th century, an age of great marine and terrestrial exploration,
Ferdinand Magellan led the first expedition to sail around the world. As a young
Portuguese noble, he served the king of Portugal, but he became involved in the
quagmire of political intrigue at court and lost the king's favor. After he
was dismissed from service by the king of Portugal, he offered to serve the
future Emperor Charles V of Spain.
```

```
A papal decree of 1493 had assigned all land in the New World west of 50 degrees
W longitude to Spain and all the land east of that line to Portugal. Magellan
offered to prove that the East Indies fell under Spanish authority. On
September 20, 1519, Magellan set sail from Spain with five ships. More than a
year later, one of these ships was exploring the topography of South America
in search of a water route across the continent. This ship sank, but the
remaining four ships searched along the southern peninsula of South America.
```

```
Finally they found the passage they sought near 50 degrees S latitude.
Magellan named this passage the Strait of All Saints, but today it is known
as the Strait of Magellan.
```

```
One ship deserted while in this passage and returned to Spain, so fewer sailors
were privileged to gaze at that first panorama of the Pacific Ocean. Those who
remained crossed the meridian now known as the International Date Line in the
early spring of 1521 after 98 days on the Pacific Ocean. During those long
days at sea, many of Magellan's men died of starvation and disease.
```

```
Later, Magellan became involved in an insular conflict in the Philippines and
was killed in a tribal battle. Only one ship and 17 sailors under the command
of the Basque navigator Elcano survived to complete the westward journey to
Spain and thus prove once and for all that the world is round, with no
precipice at the edge.'''
```

```
result = s(t)
result[0].get('summary_text')
```

```
' Ferdinand Magellan led the first expedition to sail around the world in the 16th century . Magellan set sail from Spain with five shi
ps on September 20, 1519 . He crossed the meridian now known as the International Date Line in the early spring of 1521 . Only one ship
and 17 sailors under the command of the Basque navigator Elcano survived . '
```

```
u = pipeline('fill-mask')
```

➡ No model was supplied, defaulted to distilbert/distilroberta-base and revision fb53ab8 (<https://huggingface.co/distilbert/distilroberta-base>). Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 480/480 [00:00<00:00, 40.5kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100% 331M/331M [00:01<00:00, 242MB/s]

Some weights of the model checkpoint at distilbert/distilroberta-base were not used when initializing RobertaForMaskedLM: ['roberta.pool']
- This IS expected if you are initializing RobertaForMaskedLM from the checkpoint of a model trained on another task or with another architecture
- This IS NOT expected if you are initializing RobertaForMaskedLM from the checkpoint of a model that you expect to be exactly identical

tokenizer_config.json: 100% 25.0/25.0 [00:00<00:00, 2.19kB/s]

vocab.json: 100% 899k/899k [00:00<00:00, 14.4MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 23.2MB/s]

tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 36.9MB/s]

Device set to use cpu

```
result = u("I'm teaching <mask> the board")
result
```

➡

```
[{'score': 0.4070245325565338,
  'token': 420,
  'token_str': ' across',
  'sequence': "I'm teaching across the board"},
 {'score': 0.2445068210363388,
  'token': 15,
  'token_str': ' on',
  'sequence': "I'm teaching on the board"},
 {'score': 0.13470087945461273,
  'token': 160,
  'token_str': ' off',
  'sequence': "I'm teaching off the board"},
 {'score': 0.04219989851117134,
  'token': 23,
  'token_str': ' at',
  'sequence': "I'm teaching at the board"},
 {'score': 0.02144746668636799,
  'token': 13,
  'token_str': ' for',
  'sequence': "I'm teaching for the board"}]
```

Start coding or [generate](#) with AI.