

```
!pip install transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.11/dist-packages (4.51.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from transformers) (3.18.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.30.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.31.2)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from transformers) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from transformers) (6.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from transformers) (2.32.3)
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.21.1)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11/dist-packages (from transformers) (0.5.3)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.11/dist-packages (from transformers) (4.67.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0->transformers) (2024.11.6)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub<1.0,>=0.30.0->transformers) (4.6.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.4.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->transformers) (2025.4.26)
```

```
from transformers import pipeline
```

```
a = pipeline("sentiment-analysis")
```

```
No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-2-english and revision 714eb0f (https://huggingface.co/distilbert/distilbert-base-uncased-finetuned-sst-2-english).
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 629/629 [00:00<00:00, 10.2kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, please install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.
model.safetensors: 100% 268M/268M [00:01<00:00, 191MB/s]

tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 2.70kB/s]

vocab.txt: 100% 232k/232k [00:00<00:00, 4.16MB/s]

Device set to use cpu
```

```
text = input("Enter the text:")
print(a(text))
```

```
Enter the text:i dont like this
[{'label': 'NEGATIVE', 'score': 0.8758231401443481}]
```

```
gen = pipeline('text-generation')
```

```
No model was supplied, defaulted to openai-community/gpt2 and revision 607a30d (https://huggingface.co/openai-community/gpt2).
Using a pipeline without specifying a model name and revision in production is not recommended.

config.json: 100% 665/665 [00:00<00:00, 34.1kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, please install the 'hf_xet' package.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.
model.safetensors: 100% 548M/548M [00:06<00:00, 170MB/s]

generation_config.json: 100% 124/124 [00:00<00:00, 7.34kB/s]

tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 2.01kB/s]

vocab.json: 100% 1.04M/1.04M [00:00<00:00, 28.7MB/s]

merges.txt: 100% 456k/456k [00:00<00:00, 26.1MB/s]

tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 18.8MB/s]

Device set to use cpu
```


```
text = input("Enter the text:")
t = gen(text, max_length = 200)
```

```
Enter the text:once upon a time there lived a lion in the forest
Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate inputs to the maximum length specified by the model.
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.
```

```
print(t[0].get('generated_text'))
```

 None

```
n = pipeline('ner')
```

 No model was supplied, defaulted to dbmdz/bert-large-cased-finetuned-conll03-english and revision 4c53496 (<https://huggingface.co/dbmdz/bert-large-cased-finetuned-conll03-english>)  
Using a pipeline without specifying a model name and revision in production is not recommended.

```
config.json: 100% 998/998 [00:00<00:00, 37.0kB/s]
```

Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the 'hf\_xet' package.

WARNING:huggingface\_hub.file\_download:Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download.

```
model.safetensors: 100% 1.33G/1.33G [00:18<00:00, 28.8MB/s]
```


Some weights of the model checkpoint at dbmdz/bert-large-cased-finetuned-conll03-english were not used when initializing BertForTokenClassification. This is expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task or with another architecture. This is NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be exactly identical to.

```
tokenizer_config.json: 100% 60.0/60.0 [00:00<00:00, 3.46kB/s]
```

```
vocab.txt: 100% 213k/213k [00:00<00:00, 3.65MB/s]
```


```
Device set to use cpu
```

```
text = input("Enter the text:")
result = n(text)
(result)
```

 Enter the text:i am agood girl my name is deepika from mangalagiri

```
[{'entity': 'I-PER',
  'score': np.float32(0.8506552),
  'index': 9,
  'word': 'deep',
  'start': 27,
  'end': 31},
 {'entity': 'I-PER',
  'score': np.float32(0.8543419),
  'index': 10,
  'word': '##ika',
  'start': 31,
  'end': 34},
 {'entity': 'I-LOC',
  'score': np.float32(0.49937516),
  'index': 12,
  'word': 'manga',
  'start': 40,
  'end': 45},
 {'entity': 'I-LOC',
  'score': np.float32(0.74930286),
  'index': 13,
  'word': '##lag',
  'start': 45,
  'end': 48},
 {'entity': 'I-LOC',
  'score': np.float32(0.60719955),
  'index': 14,
  'word': '##iri',
  'start': 48,
  'end': 51}]
```

```
qa = pipeline('question-answering')
```

 No model was supplied, defaulted to distilbert/distilbert-base-cased-distilled-squad and revision 564e9b5 (<https://huggingface.co/distilbert/distilbert-base-cased-distilled-squad>)  
Using a pipeline without specifying a model name and revision in production is not recommended.

```
config.json: 100% 473/473 [00:00<00:00, 19.4kB/s]
```

Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the 'hf\_xet' package.

WARNING:huggingface\_hub.file\_download:Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download.

```
model.safetensors: 100% 261M/261M [00:08<00:00, 13.3MB/s]
```

```
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 2.27kB/s]
```

```
vocab.txt: 100% 213k/213k [00:00<00:00, 3.93MB/s]
```

```
tokenizer.json: 100% 436k/436k [00:00<00:00, 14.3MB/s]
```

```
Device set to use cpu
```

```
c = '''Bity bought some butter,but the butter was bitter butter.
So Bity bought some better butter to make the bitter butter better.'''
```

```
q = "What does bitty bought?"
```

```
result = qa({'question': q, 'context': c})
print(result)
```

```
➦ {'score': 0.5355311036109924, 'start': 12, 'end': 23, 'answer': 'some butter'}
/usr/local/lib/python3.11/dist-packages/transformers/pipelines/question_answering.py:390: FutureWarning: Passing a list of SQuAD example warnings.warn()
```

```
s = pipeline('summarization')
```

```
➦ No model was supplied, defaulted to sshleifer/distilbart-cnn-12-6 and revision a4f8f3e (https://huggingface.co/sshleifer/distilbart-cnn-12-6).
Using a pipeline without specifying a model name and revision in production is not recommended.
```

```
config.json: 100% 1.80k/1.80k [00:00<00:00, 85.6kB/s]
pytorch_model.bin: 100% 1.22G/1.22G [00:25<00:00, 30.8MB/s]
model.safetensors: 100% 1.22G/1.22G [00:20<00:00, 18.3MB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 1.09kB/s]
vocab.json: 100% 899k/899k [00:00<00:00, 9.11MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 6.23MB/s]
Device set to use cpu
```

```
t = '''Bity bought some butter,but the butter was bitter butter.
So Bity bought some better butter to make the bitter butter better.'''
```

```
result = s(t)
print(result[0].get('summary_text'))
```

```
➦ Your max_length is set to 142, but your input_length is only 30. Since this is a summarization task, where outputs shorter than the input
Bity bought some butter,but the butter was bitter butter . So she bought some better butter to make the bitter butter better . Bity buy
```

```
u = pipeline('fill-mask')
```

```
➦ No model was supplied, defaulted to distilbert/distilroberta-base and revision fb53ab8 (https://huggingface.co/distilbert/distilroberta-base).
Using a pipeline without specifying a model name and revision in production is not recommended.
```

```
config.json: 100% 480/480 [00:00<00:00, 27.3kB/s]
Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download
model.safetensors: 100% 331M/331M [00:08<00:00, 38.5MB/s]
Some weights of the model checkpoint at distilbert/distilroberta-base were not used when initializing RobertaForMaskedLM: ['roberta.pool']
- This IS expected if you are initializing RobertaForMaskedLM from the checkpoint of a model trained on another task or with another architecture
- This IS NOT expected if you are initializing RobertaForMaskedLM from the checkpoint of a model that you expect to be exactly identical
tokenizer_config.json: 100% 25.0/25.0 [00:00<00:00, 1.93kB/s]
vocab.json: 100% 899k/899k [00:00<00:00, 10.3MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 11.9MB/s]
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 23.6MB/s]
Device set to use cpu
```

```
result = u('Bity bought some <mask> butter')
(result)
```

```
➦ [{'score': 0.646386981010437,
  'token': 22189,
  'token_str': ' peanut',
  'sequence': 'Bity bought some peanut butter'},
 {'score': 0.0758957713842392,
  'token': 32473,
  'token_str': ' almond',
  'sequence': 'Bity bought some almond butter'},
 {'score': 0.02743339166045189,
  'token': 15848,
  'token_str': ' vegan',
  'sequence': 'Bity bought some vegan butter'},
 {'score': 0.01600140519440174,
```

```
'token': 24136,  
'token_str': ' melted',  
'sequence': 'Bity bought some melted butter'},  
{'score': 0.015843648463487625,  
'token': 17145,  
'token_str': ' nut',  
'sequence': 'Bity bought some nut butter'}}
```

Start coding or [generate](#) with AI.